

Alex Chen

Title: Senior AI Software Engineer

Email: alex.chen@email.com

Professional Summary:

Full-stack AI engineer specializing in local LLM deployment and custom fine-tuning. Passionate about building privacy-first GenAI applications that operate completely on edge hardware without relying on OpenAI.

Experience:

Lead AI Engineer | Vertex Solutions (2023 - Present)

- Architected a local enterprise RAG system using Next.js for the frontend and a FastAPI backend to orchestrate document ingestion.
- Fine-tuned LLaMA 3 8B using Unsloth and QLoRA to generate strict JSON schemas from unstructured medical PDFs.
- Deployed quantized GGUF models on constrained local GPU hardware (RTX 3090/3060) utilizing Ollama and vLLM to strictly manage VRAM utilization and KV caching.
- Built asynchronous streaming text UIs using React and Tailwind CSS.

Backend Engineer | DataFlow Tech (2020 - 2023)

- Developed high-throughput REST APIs using Python and FastAPI.
- Implemented sequential processing queues to handle heavy PDF parsing pipelines using PyMuPDF.

Skills:

Python, TypeScript, React, Next.js, FastAPI, LLaMA 3, Unsloth, QLoRA, vLLM, Ollama, llama.cpp, Tailwind CSS, PyTorch.